

Lab Sheet 4: Basics of data analysis in R and a Simple Sentiment Analysis

Lab Sheet 4a: Basics of data analysis in R

Introduction

In this section, we will explore basic data transformation techniques using the `dplyr` package with the `nycflights13` dataset, which contains information about all flights that departed from NYC in 2013.

Load Necessary Packages

```
# Load required packages
library(nycflights13) # Dataset containing flight data
library(tidyverse)    # A collection of R packages for data manipulation and visualization
```

Data Transformation with dplyr

1. Filtering Data We can use the `filter()` function to subset our data based on certain conditions.

```
# Filter flights that had a departure delay greater than 120 minutes
flights |>
  filter(dep_delay > 120)
```

```
## # A tibble: 9,723 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>   <int>
## 1  2013     1     1     848           1835         853    1001
## 2  2013     1     1     957           733         144    1056
## 3  2013     1     1    1114           900         134    1447
## 4  2013     1     1    1540          1338         122    2020
## 5  2013     1     1    1815          1325         290    2120
## 6  2013     1     1    1842          1422         260    1958
## 7  2013     1     1    1856          1645         131    2212
## 8  2013     1     1    1934          1725         129    2126
## 9  2013     1     1    1938          1703         155    2109
## 10 2013     1     1    1942          1705         157    2124
## # i 9,713 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Flights that departed on January 1
flights |>
  filter(month == 1 & day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>   <int>
```

```
## 1 2013 1 1 517 515 2 830
## 2 2013 1 1 533 529 4 850
## 3 2013 1 1 542 540 2 923
## 4 2013 1 1 544 545 -1 1004
## 5 2013 1 1 554 600 -6 812
## 6 2013 1 1 554 558 -4 740
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # i 832 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## # dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## # minute <dbl>, time_hour <dtm>
```

```
# Flights that departed in November or December
flights |>
  filter(month == 11 | month == 12)
```

```
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013    11     1       5          2359         6       352
## 2 2013    11     1      35          2250        105      123
## 3 2013    11     1     455           500         -5      641
## 4 2013    11     1     539           545         -6      856
## 5 2013    11     1     542           545         -3      831
## 6 2013    11     1     549           600        -11      912
## 7 2013    11     1     550           600        -10      705
## 8 2013    11     1     554           600         -6      659
## 9 2013    11     1     554           600         -6      826
## 10 2013    11     1     554           600         -6      749
## # i 55,393 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## # dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## # minute <dbl>, time_hour <dtm>
```

```
# Another way to filter for November or December
flights |>
  filter(month %in% c(11, 12))
```

```
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013    11     1       5          2359         6       352
## 2 2013    11     1      35          2250        105      123
## 3 2013    11     1     455           500         -5      641
## 4 2013    11     1     539           545         -6      856
## 5 2013    11     1     542           545         -3      831
## 6 2013    11     1     549           600        -11      912
## 7 2013    11     1     550           600        -10      705
## 8 2013    11     1     554           600         -6      659
## 9 2013    11     1     554           600         -6      826
## 10 2013    11     1     554           600         -6      749
```

```
## # i 55,393 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. Arranging Data The `arrange()` function helps us sort the data.

```
# Arrange flights by year, month, day, and departure time
flights |>
  arrange(year, month, day, dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # i 336,766 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Arrange flights by departure delay in descending order
flights |>
  arrange(desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641             900        1301    1242
## 2  2013     6    15    1432            1935        1137    1607
## 3  2013     1    10    1121            1635        1126    1239
## 4  2013     9    20    1139            1845        1014    1457
## 5  2013     7    22     845            1600        1005    1044
## 6  2013     4    10    1100            1900         960    1342
## 7  2013     3    17    2321             810         911     135
## 8  2013     6    27     959            1900         899    1236
## 9  2013     7    22    2257             759         898     121
## 10 2013    12     5     756            1700         896    1058
## # i 336,766 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. Distinct Values To find unique values in our dataset, we can use `distinct()`.

```
# Remove duplicate rows from the dataset
flights |>
  distinct()
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # i 336,766 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Keep only unique combinations of origin and destination
flights |>
  distinct(origin, dest, .keep_all = TRUE)
```

```
## # A tibble: 224 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # i 214 more rows
## # i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Count unique combinations of origin and destination
flights |>
  count(origin, dest, sort = TRUE)
```

```
## # A tibble: 224 x 3
##   origin dest      n
##   <chr>  <chr> <int>
## 1 JFK    LAX   11262
## 2 LGA    ATL   10263
```

```
## 3 LGA ORD 8857
## 4 JFK SFO 8204
## 5 LGA CLT 6168
## 6 EWR ORD 6100
## 7 JFK BOS 5898
## 8 LGA MIA 5781
## 9 JFK MCO 5464
## 10 EWR BOS 5327
## # i 214 more rows
```

4. **Creating New Variables with `mutate()`** The `mutate()` function allows us to create new variables.

```
# Calculate the gain in minutes and speed of flights
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed = distance / air_time * 60 # Speed in miles per hour
  )
```

```
## # A tibble: 336,776 x 21
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517             515         2     830
## 2  2013     1     1     533             529         4     850
## 3  2013     1     1     542             540         2     923
## 4  2013     1     1     544             545        -1    1004
## 5  2013     1     1     554             600        -6     812
## 6  2013     1     1     554             558        -4     740
## 7  2013     1     1     555             600        -5     913
## 8  2013     1     1     557             600        -3     709
## 9  2013     1     1     557             600        -3     838
## 10 2013     1     1     558             600        -2     753
## # i 336,766 more rows
## # i 14 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, gain <dbl>, speed <dbl>
```

5. **Using Pipes** The pipe operator (`|>`) lets us chain commands together for clearer and more concise code.

```
# Filter for flights to IAH, calculate speed, and arrange by speed
flights |>
  filter(dest == "IAH") |>
  mutate(speed = distance / air_time * 60) |>
  select(year:day, dep_time, carrier, flight, speed) |>
  arrange(desc(speed))
```

```
## # A tibble: 7,198 x 7
##   year month   day dep_time carrier flight speed
##   <int> <int> <int>   <int> <chr>    <int> <dbl>
## 1  2013     7     9     707 UA      226 522.
## 2  2013     8    27    1850 UA     1128 521.
## 3  2013     8    28     902 UA     1711 519.
## 4  2013     8    28    2122 UA     1022 519.
## 5  2013     6    11    1628 UA     1178 515.
```

```
## 6 2013      8    27      1017 UA          333 515.
## 7 2013      8    27      1205 UA          1421 515.
## 8 2013      8    27      1758 UA           302 515.
## 9 2013      9    27        521 UA           252 515.
## 10 2013     8    28        625 UA           559 515.
## # i 7,188 more rows
```

6. Grouping Data Grouping data allows us to perform calculations on subsets of the data.

```
# Group by month and calculate average departure delay
flights |>
  group_by(month) |>
  summarize(
    avg_delay = mean(dep_delay, na.rm = TRUE) # Handle missing values
  )
```

```
## # A tibble: 12 x 2
##   month avg_delay
##   <int>   <dbl>
## 1     1    10.0
## 2     2    10.8
## 3     3    13.2
## 4     4    13.9
## 5     5    13.0
## 6     6    20.8
## 7     7    21.7
## 8     8    12.6
## 9     9     6.72
## 10    10     6.24
## 11    11     5.44
## 12    12    16.6
```

Conclusion

In this section, we've covered essential data transformation techniques using `dplyr`. These skills are foundational for data analysis in R.

Lab Sheet 4b: A Simple Sentiment Analysis

Introduction

In this section, we will perform sentiment analysis on a news article using the `tidytext` and `rvest` packages. We will scrape the article, process the text, and visualize the results.

Load Necessary Packages

```
# Load required packages for sentiment analysis
library(tidytext) # Text mining
library(rvest)    # Web scraping
library(textdata) # Sentiment datasets
library(wordcloud) # Word cloud visualization
library(RColorBrewer) # Color palettes for visualizations
library(wordcloud2) # Interactive word clouds
```

Scrape and Extract Text from a Website

We will use the `rvest` package to scrape a news article.

```
# Define the URL of the article
url <- 'https://indianexpress.com/article/technology/science/nobel-prize-physics-john-hopfield-geoffrey'

# Read the HTML content from the webpage
news <- read_html(url)
```

Extract Text

We will extract all the paragraphs from the article and store them in a tibble.

```
# Scrape and extract all paragraphs (<p>) and store them in a tibble
text <- tibble(
  news %>%
    html_elements('p') %>% # Select all paragraph elements (<p>) from the HTML content
    html_text() # Extract text from the paragraph elements
) %>%
  rename('text' = 1) # Rename the column to "text"
```

Get Sentiments

We will load the NRC sentiment dataset, which categorizes words into different emotions.

```
# Load sentiment data
sentiments <- get_sentiments('nrc')
```

Tokenize the Text

We will split the text into individual words for analysis.

```
# Separate all texts into individual words
tokens <- text %>%
  unnest_tokens(input = text, output = word) %>%
  filter(!grepl('[0-9]', word)) # Remove numbers
```

Remove Stop-Words and Count Frequency

We will remove common stop-words and count the frequency of each word.

```
# Remove stop-words and count the frequency of each word
word_freq <- tokens %>%
  anti_join(stop_words) %>%
  count(word, sort=TRUE)
```

```
## Joining with `by = join_by(word)`
```

Visualize Word Frequencies with Word Clouds

We can visualize the most frequent words using word clouds.

```
# Set seed for reproducibility
set.seed(1234)

# Basic word cloud
wordcloud(words = word_freq$word,
          freq = word_freq$n,
```

[illegible]

```
## PhantomJS not found. You can install it with webshot::install_phantomjs(). If it is installed, please
```


Next, we will associate each word with its sentiment score.

8

Plot the Sentiment Frequencies

Finally, we will visualize the sentiment frequencies using a bar plot.

```
# Create a bar plot of sentiment frequencies
gg <- freq_count %>%
  ggplot(aes(x= sentiment, y= n, fill= sentiment)) +
  geom_col(show.legend = FALSE) +
  geom_text(aes(label=n), size=5, position = position_stack(vjust = 0.5)) + # Center labels
  labs(x= 'Sentiment', y= 'Frequency') +
  ggtitle('Sentiment Frequencies from Article') +
  theme_minimal() # Clean theme
gg
```



Conclusion

In this section, we performed sentiment analysis on a news article, demonstrating how to scrape text data, analyze it for sentiment, and visualize the results.

Additional Notes

- Ensure you have all the necessary packages installed before running the scripts.
- Explore additional datasets and articles for further practice with sentiment analysis.